




Making servers gossip


Tomás Fernandes


“UWCS is giving free pizza
during Friday Night Gaming”


- Some Computer Science student, circa 2022

 Freshers











UWCS goes bankrupt!

UWCS goes bankrupt!


So what?

Gossip protocol


Inspired by **human gossip** and **epidemics**



Overlay / Topology




Peer




Round 1


Fanout = 3




Round 1




Round 1




Round 2




Round 2




Round 2




Round 2



Round 2



Round 2



Scalable!

$$O(\log n) = \sim 1.89$$

rounds to spread the rumor with everyone





We gossip with a **fixed** number of people



No recovery action is taken

Fast convergence
& eventually consistent








$$2^n$$


The free pizza rumor spreads exponentially across the Warwick campus


Decentralized














Running the same software, no central coordinators


Fault Tolerant










Robust*







Amazon S3



Amazon S3

Scalable service for object-storage offered by AWS


Uses a Gossip protocol to spread the state of each server


2008 Loss of availability


Messages with a corrupted bit were circulated


Amazon had to shut down communication between S3 servers


The state was manually fixed to restore the service











Byzantine Faults

Don't bring a system to a full-stop

Maybe be caused by malicious or flaky peers

How do we know who to
speak with?

Keep track of every
peer?

Keep track of every
peer?


$$O(\textit{horrible}) = O(n)$$


Not scalable!

HyParView

A membership protocol for reliable gossip-based broadcast (Joao et al.)

Hybrid **Partial View**







$$O(\log n)$$


Offers strong resilience to failures of peers
(as high as 90%)


Takes advantage of TCP

Overlay / Topology









Applications

Cluster Membership (Cassandra)

Failure Detection (HashiCorp Serf)

Distributed training of ML models

Compute aggregates

Disseminate metadata (CockroachDB)

Resources

[Simulator](#)

[Convergence Simulator](#)

[Gossip Dissemination – Martin Fowler](#)

[Introduction to Gossip – Felix Lopez](#)

[HyParView](#)

[Amazon S3 Outage](#) (Web Archive)

[HyParView: F#](#)

[HyParView: Rust](#)

[The Peer Sampling Service: Experimental Evaluation of Unstructured Gossip-Based Implementations?](#)



Q&A

tomff.com/talks/gossip.pdf